

duck_db_exploration

December 27, 2024

1 Duck DB Exploration with FitBit Data

The purpose of this notebook is to explore the versatility of [DuckDB](#) while exploring my data from my Fitbit. If you [download your data from fitbit](#) you will end up with a zipped file with many directories and a mix of JSON files and CSV files. Building a custom parser to get this data and organize it properly can be time consuming, especially since the JSON files have some deep nesting. Instead what we can do is have DuckDB read in these JSON files, execute SQL against them and then merge data together to explore it. My end goal for this exploration is to create a model to predict my HRV, so what I am going to need is data on my HRV, and I'm also going to bring in my sleep data and my daily heart rate data to see if there are good predictors there. I am fully aware that the HRV is likely not accurate on my Fitbit, but let's have some fun anyway! I will not be building models or building visualizations in this notebook, I will save that for another notebook.

1.1 Getting Started

Before we get going, just a little explanation of the data. Different versions of Fitbits will likely have different sensors and therefore different data. I haven't done any research to know if the directories are the same and I am not really concerned about that. I have a Charge 5 and I have noticed that there is a TON more data here than the application shows you. We have a lots of directories and data, let's take a look at how many files we have and what types. **Note: when I was developing this I used the `.show()`** function for displaying results directly from DuckDB.** This did a really good job displaying nested data in a tabular format, but the display looked terrible on the HTML export of the Jupyter Notebook, so I converted everything to a dataframe and displayed that instead, but the show function works just fine in jupyter, just not the export.**

```
[1]: import pandas as pd
import os
from collections import defaultdict

def list_directories_and_count_files_by_type_with_summary(root_dir):
    overall_file_type_counts = defaultdict(int) # To store the overall summary

    for dirpath, dirnames, filenames in os.walk(root_dir):
        file_type_counts = defaultdict(int)

        for filename in filenames:
            file_extension = os.path.splitext(filename)[1].lower() # Get the
↪file extension
```

```

        file_type_counts[file_extension] += 1
        overall_file_type_counts[file_extension] += 1 # Add to overall
    summary

    print(f"Directory: {dirpath}")
    for file_type, count in file_type_counts.items():
        print(f" {file_type} if file_type else '[No Extension]': {count}")
    print("-" * 40)

    # Print the overall summary
    print("Summary of all file types across all directories:")
    for file_type, total_count in overall_file_type_counts.items():
        print(f" {file_type} if file_type else '[No Extension]':
    {total_count}")
    print("-" * 40)

root_directory = './data/unzipped/Takeout'
list_directories_and_count_files_by_type_with_summary(root_directory)

```

```

Directory: ./data/unzipped/Takeout
.html: 1
-----
Directory: ./data/unzipped/Takeout\Fitbit
-----
Directory: ./data/unzipped/Takeout\Fitbit\Account Changes
.csv: 1
-----
Directory: ./data/unzipped/Takeout\Fitbit\Active Zone Minutes (AZM)
.csv: 16
-----
Directory: ./data/unzipped/Takeout\Fitbit\Activity Goals
.txt: 1
.csv: 1
-----
Directory: ./data/unzipped/Takeout\Fitbit\Atrial Fibrillation ECG
.csv: 6
.txt: 1
-----
Directory: ./data/unzipped/Takeout\Fitbit\Biometrics
.txt: 1
.csv: 210
-----
Directory: ./data/unzipped/Takeout\Fitbit\Daily Readiness
.txt: 1
.csv: 15
-----
Directory: ./data/unzipped/Takeout\Fitbit\Discover

```

```

.txt: 1
-----
Directory: ./data/unzipped/Takeout\Fitbit\Fitbit Care or Programs
.txt: 1
-----
Directory: ./data/unzipped/Takeout\Fitbit\Fitbit Friends
.txt: 1
.csv: 1
-----
Directory: ./data/unzipped/Takeout\Fitbit\Fitbit Premium
.txt: 1
.csv: 2
-----
Directory: ./data/unzipped/Takeout\Fitbit\Global Export Data
.json: 1116
.csv: 473
-----
Directory: ./data/unzipped/Takeout\Fitbit\Guided Programs
.txt: 1
-----
Directory: ./data/unzipped/Takeout\Fitbit\Health Fitness Data_GoogleData
.txt: 5
.csv: 6
-----
Directory: ./data/unzipped/Takeout\Fitbit\Heart Rate
.csv: 2
.txt: 1
-----
Directory: ./data/unzipped/Takeout\Fitbit\Heart Rate Variability
.csv: 1430
.txt: 4
-----
Directory: ./data/unzipped/Takeout\Fitbit\Menstrual Health
.txt: 1
.csv: 4
-----
Directory: ./data/unzipped/Takeout\Fitbit\Mindfulness
.csv: 3
.txt: 1
-----
Directory: ./data/unzipped/Takeout\Fitbit\Oxygen Saturation (SpO2)
.csv: 478
.txt: 2
-----
Directory: ./data/unzipped/Takeout\Fitbit\Paired Devices
.txt: 1
.csv: 5
-----

```

Directory: ./data/unzipped/Takeout\Fitbit\Physical Activity_GoogleData
 .csv: 574
 .txt: 9

Directory: ./data/unzipped/Takeout\Fitbit\Sleep
 .txt: 1
 .csv: 1

Directory: ./data/unzipped/Takeout\Fitbit\Sleep Score
 .csv: 1

Directory: ./data/unzipped/Takeout\Fitbit\Snore and Noise Detect
 .txt: 1

Directory: ./data/unzipped/Takeout\Fitbit\Social
 .csv: 1
 .txt: 1
 .png: 1

Directory: ./data/unzipped/Takeout\Fitbit\Stress Journal
 .txt: 1

Directory: ./data/unzipped/Takeout\Fitbit\Stress Score
 .txt: 1
 .csv: 1

Directory: ./data/unzipped/Takeout\Fitbit\Temperature
 .csv: 490

Directory: ./data/unzipped/Takeout\Fitbit\Transactions
 .txt: 1

Directory: ./data/unzipped/Takeout\Fitbit\User Security Data
 .csv: 2

Directory: ./data/unzipped/Takeout\Fitbit\Your Profile
 .jpg: 1
 .csv: 1
 .txt: 1

Summary of all file types across all directories:

 .html: 1
 .csv: 3724
 .txt: 40
 .json: 1116
 .png: 1
 .jpg: 1

That is a lot of files, they're not very big but each style of file likely has it's own way of storing stuff. If it's not a CSV it's going to be annoying to get what we need. In addition to that, if we wanted to write SQL against it, loading data into a database like PostgreSQL would be annoyingly time consuming to setup, especially if our goal is to explore and analyze the data and not necessarily create an application out of it. We just want to organize the data well using simple SQL syntax and then get what we need for our analysis and that's what **DuckDB** is going to help us with. Okay enough chatting let's do something.

1.2 DuckDB - Getting Started

These few lines of code are all we need to load our sleep data. First we import the library, then we make a connection, you don't have to supply a string argument which creates a permanent database. If you don't supply an argument it's like having a temp database, all the rest of the commands in the notebook will work. I did it because later, I will want to store some data in a table, but for now it's not necessary. Finally the 3rd line is where we get a lot of help, we are going to read in all of the JSON files in the *Global Export Data* directory that begin with "sleep-". **DuckDB** will read the json data and store it in the `sleep_data` relation (what **DuckDB** uses).

```
[2]: import duckdb
conn = duckdb.connect('./data/fitbit_db.duckdb') # creates a database
sleep_data = duckdb.read_json("./data/unzipped/Takeout/Fitbit/Global Export_
↳Data/sleep-*.json")
print(f'The object type for sleep_data is {type(sleep_data)}')
```

The object type for `sleep_data` is `<class 'duckdb.duckdb.DuckDBPyRelation'>`

This did a pretty good job of reading in the data, it parsed the **dateOfSleep** properly and got the **mainSleep** as a *boolean* instead of a *varchar* which is pretty nice. Some more interesting points is that the **levels** field has a datatype of *struct* and inside we see all the nested JSON data. The first section is called **summary** which holds aggregate information about my different stages of sleep (deep, wake, light, & rem). It also has a bunch of NULL features tacked on to the end. After **summary** we have an array of **data** which shows the chronological sequence of each sleep stage and it's duration in seconds. There's also an array of data called **shortData** which contains small blips of time in chronological order. That makes sense because when I look at the *Sleep Timeline* section on the Fitbit app there are often short small bursts of time in the Awake section. This is likely the data that is displayed there.

Let's reimport the data and convert those timestamps so that we can use this data properly.

```
[3]: sleep_data = duckdb.read_json("./data/unzipped/Takeout/Fitbit/Global Export_
↳Data/sleep-*.json", timestamp_format="%Y-%m-%dT%H:%M:%S.%g")
display(sleep_data.df().head(5))
```

	logId	dateOfSleep	startTime	endTime	duration	\
0	42724299942	2023-09-10	2023-09-09 21:34:00	2023-09-10 06:11:00	31020000	
1	42712673606	2023-09-09	2023-09-08 22:14:00	2023-09-09 05:37:30	26580000	
2	42703109427	2023-09-08	2023-09-07 21:48:00	2023-09-08 05:23:30	27300000	
3	42688951099	2023-09-07	2023-09-06 21:58:00	2023-09-07 05:46:30	28080000	
4	42676589452	2023-09-06	2023-09-05 22:10:30	2023-09-06 05:37:00	26760000	

	minutesToFallAsleep	minutesAsleep	minutesAwake	minutesAfterWakeup	\
0	0	457	60	0	
1	0	400	43	0	
2	0	398	57	0	
3	0	413	55	1	
4	0	376	70	10	

	timeInBed	efficiency	type	infoCode	logType	\
0	517	98	stages	0	auto_detected	
1	443	100	stages	0	auto_detected	
2	455	99	stages	0	auto_detected	
3	468	98	stages	0	auto_detected	
4	446	96	stages	0	auto_detected	

	levels	mainSleep
0	{'summary': {'deep': {'count': 5, 'minutes': 1...	True
1	{'summary': {'deep': {'count': 4, 'minutes': 6...	True
2	{'summary': {'deep': {'count': 6, 'minutes': 7...	True
3	{'summary': {'deep': {'count': 5, 'minutes': 8...	True
4	{'summary': {'deep': {'count': 3, 'minutes': 8...	True

That was really easy and fast, what I'd like to do now is explore this dataset a bit. The goal here is to show simple ways to access and point out some interesting features. First you can describe the data, similar to what you are able to do with pandas.

```
[4]: sleep_data.describe().df()
```

```
[4]:
```

	aggr	logId	dateOfSleep	startTime	endTime	\
0	count	4.810000e+02	481	481	481	
1	mean	4.506723e+10	None	None	None	
2	stddev	1.513614e+09	None	None	None	
3	min	4.240092e+10	2023-08-13	2023-08-12 22:16:00	2023-08-13 05:54:30	
4	max	4.759902e+10	2024-11-27	2024-11-26 22:39:00	2024-11-27 04:01:30	
5	median	4.511957e+10	None	None	None	

	duration	minutesToFallAsleep	minutesAsleep	minutesAwake	\
0	4.810000e+02	481.0	481.000000	481.000000	
1	2.732532e+07	0.0	406.800416	48.496881	
2	3.387921e+06	0.0	52.882236	11.543419	
3	5.460000e+06	0.0	82.000000	1.000000	
4	3.558000e+07	0.0	543.000000	90.000000	
5	2.748000e+07	0.0	410.000000	48.000000	

	minutesAfterWakeup	timeInBed	efficiency	type	infoCode	\
0	481.000000	481.000000	481.000000	481	481.000000	
1	0.557173	455.422037	98.168399	None	0.008316	
2	1.348361	56.465353	1.677895	None	0.128831	
3	0.000000	91.000000	88.000000	classic	0.000000	

4	10.000000	593.000000	100.000000	stages	2.000000
5	0.000000	458.000000	99.000000	None	0.000000

	logType	levels	mainSleep
0	481	481	481
1	None	None	None
2	None	None	None
3	auto_detected	{'summary': {'deep': {'count': 1, 'minutes': 4...	false
4	auto_detected	{'summary': {'deep': NULL, 'wake': NULL, 'ligh...	true
5	None	None	None

I can see that I have 481 entries in this dataset, but from the **dateOfSleep** data I'm not convinced I have 1 record for every night as the date difference between 2023-08-13 and 2024-11-27 is 471 days not 481. I also know that some nights I probably forgot to wear my fitbit so I know the number should be even less than 471. I can also see that some columns aren't going to be very useful to me, like **minutesToFallAsleep** has no mean or stddev so it never changes from 0, similarly **logType** has the same 'auto_detected' value. Let's write some SQL to take a look at the data and see what's up.

```
[5]: display(duckdb.sql("SELECT dateOfSleep, count(*) as count FROM sleep_data GROUP_
    ↳BY dateOfSleep Having count(*) > 1").df())
display(duckdb.sql("select count(distinct dateOfSleep) as distinct_dates from_
    ↳sleep_data").df())
```

	dateOfSleep	count
0	2023-11-09	2
1	2024-07-06	2
2	2023-10-10	2
3	2023-09-10	2
4	2024-03-08	2
5	2024-05-07	2
6	2024-08-05	2
7	2024-06-06	2
8	2024-10-04	2
9	2024-01-08	2
10	2023-12-09	2
11	2024-04-07	2
12	2024-02-07	2
13	2024-09-04	2
14	2023-11-14	2
15	2024-11-03	2

	distinct_dates
0	465

I was right 465 unique dates, but 481 rows so I have some duplicates for some days...something is funky with that data. Before I work on that I just wanted to point out what **DuckDB** was able to do for us, our 'table' is the object that we read in. We can easily write SQL syntax against files or anything we read in just by using it like we would a table in SQL...pretty cool.

Now I want to take a look at the data from those duplicate dates, SQL syntax to the rescue!

```
[6]: display(duckdb.sql("""SELECT s.* from sleep_data s where dateOfSleep in
      (SELECT dateOfSleep FROM sleep_data GROUP BY dateOfSleep Having
      ↪count(*) > 1)
      order by dateOfSleep""").df())
display(duckdb.sql("""SELECT s.logId, unnest(levels.summary.deep) from
      ↪sleep_data s where dateOfSleep in
      (SELECT dateOfSleep FROM sleep_data GROUP BY dateOfSleep Having
      ↪count(*) > 1)
      order by dateOfSleep""").df())
```

	logId	dateOfSleep		startTime		endTime		duration	\
0	42724299942	2023-09-10	2023-09-09	21:34:00	2023-09-10	06:11:00	31020000		
1	42724299942	2023-09-10	2023-09-09	21:34:00	2023-09-10	06:11:00	31020000		
2	43071124476	2023-10-10	2023-10-09	22:29:00	2023-10-10	05:40:30	25860000		
3	43071124476	2023-10-10	2023-10-09	22:29:00	2023-10-10	05:40:30	25860000		
4	43415888071	2023-11-09	2023-11-08	22:12:00	2023-11-09	05:32:30	26400000		
5	43415888071	2023-11-09	2023-11-08	22:12:00	2023-11-09	05:32:30	26400000		
6	43483210808	2023-11-14	2023-11-14	22:23:30	2023-11-14	23:55:00	5460000		
7	43472303148	2023-11-14	2023-11-13	23:19:30	2023-11-14	06:01:30	24120000		
8	43758999257	2023-12-09	2023-12-08	21:45:30	2023-12-09	04:42:00	24960000		
9	43758999257	2023-12-09	2023-12-08	21:45:30	2023-12-09	04:42:00	24960000		
10	44096210623	2024-01-08	2024-01-07	22:05:00	2024-01-08	05:31:00	26760000		
11	44096210623	2024-01-08	2024-01-07	22:05:00	2024-01-08	05:31:00	26760000		
12	44452696017	2024-02-07	2024-02-06	21:17:00	2024-02-07	05:30:30	29580000		
13	44452696017	2024-02-07	2024-02-06	21:17:00	2024-02-07	05:30:30	29580000		
14	44796156145	2024-03-08	2024-03-07	21:00:00	2024-03-08	05:31:30	30660000		
15	44796156145	2024-03-08	2024-03-07	21:00:00	2024-03-08	05:31:30	30660000		
16	45141553280	2024-04-07	2024-04-06	21:26:30	2024-04-07	06:41:30	33300000		
17	45141553280	2024-04-07	2024-04-06	21:26:30	2024-04-07	06:41:30	33300000		
18	45472556013	2024-05-07	2024-05-06	22:08:00	2024-05-07	05:21:30	25980000		
19	45472556013	2024-05-07	2024-05-06	22:08:00	2024-05-07	05:21:30	25980000		
20	45804854753	2024-06-06	2024-06-05	20:59:00	2024-06-06	05:32:30	30780000		
21	45804854753	2024-06-06	2024-06-05	20:59:00	2024-06-06	05:32:30	30780000		
22	46126383325	2024-07-06	2024-07-05	22:10:00	2024-07-06	06:11:30	28860000		
23	46126383325	2024-07-06	2024-07-05	22:10:00	2024-07-06	06:11:30	28860000		
24	46439171327	2024-08-05	2024-08-04	22:01:30	2024-08-05	05:31:30	27000000		
25	46439171327	2024-08-05	2024-08-04	22:01:30	2024-08-05	05:31:30	27000000		
26	46747290318	2024-09-04	2024-09-03	21:45:00	2024-09-04	06:04:30	29940000		
27	46747290318	2024-09-04	2024-09-03	21:45:00	2024-09-04	06:04:30	29940000		
28	47046675063	2024-10-04	2024-10-03	22:12:30	2024-10-04	05:51:00	27480000		
29	47046675063	2024-10-04	2024-10-03	22:12:30	2024-10-04	05:51:00	27480000		
30	47353917202	2024-11-03	2024-11-02	21:30:00	2024-11-03	05:48:30	29880000		
31	47353917202	2024-11-03	2024-11-02	21:30:00	2024-11-03	05:48:30	29880000		

	minutesToFallAsleep	minutesAsleep	minutesAwake	minutesAfterWakeup	\
0	0	457	60	0	

1	0	457	60	0
2	0	384	47	0
3	0	384	47	0
4	0	399	41	1
5	0	399	41	1
6	0	82	9	0
7	0	362	40	7
8	0	373	43	2
9	0	373	43	2
10	0	395	51	1
11	0	395	51	1
12	0	427	66	0
13	0	427	66	0
14	0	465	46	0
15	0	465	46	0
16	0	499	56	0
17	0	499	56	0
18	0	385	48	0
19	0	385	48	0
20	0	448	65	0
21	0	448	65	0
22	0	418	63	0
23	0	418	63	0
24	0	418	32	1
25	0	418	32	1
26	0	445	54	0
27	0	445	54	0
28	0	409	49	0
29	0	409	49	0
30	0	427	71	0
31	0	427	71	0

	timeInBed	efficiency	type	infoCode	logType \
0	517	98	stages	0	auto_detected
1	517	98	stages	0	auto_detected
2	431	99	stages	0	auto_detected
3	431	99	stages	0	auto_detected
4	440	98	stages	0	auto_detected
5	440	98	stages	0	auto_detected
6	91	90	classic	2	auto_detected
7	402	98	stages	0	auto_detected
8	416	97	stages	0	auto_detected
9	416	97	stages	0	auto_detected
10	446	98	stages	0	auto_detected
11	446	98	stages	0	auto_detected
12	493	99	stages	0	auto_detected
13	493	99	stages	0	auto_detected
14	511	100	stages	0	auto_detected

15	511	100	stages	0	auto_detected
16	555	99	stages	0	auto_detected
17	555	99	stages	0	auto_detected
18	433	98	stages	0	auto_detected
19	433	98	stages	0	auto_detected
20	513	99	stages	0	auto_detected
21	513	99	stages	0	auto_detected
22	481	98	stages	0	auto_detected
23	481	98	stages	0	auto_detected
24	450	99	stages	0	auto_detected
25	450	99	stages	0	auto_detected
26	499	100	stages	0	auto_detected
27	499	100	stages	0	auto_detected
28	458	98	stages	0	auto_detected
29	458	98	stages	0	auto_detected
30	498	98	stages	0	auto_detected
31	498	98	stages	0	auto_detected

		levels	mainSleep
0	{'summary': {'deep': {'count': 5, 'minutes': 1...		True
1	{'summary': {'deep': {'count': 5, 'minutes': 1...		True
2	{'summary': {'deep': {'count': 5, 'minutes': 8...		True
3	{'summary': {'deep': {'count': 5, 'minutes': 8...		True
4	{'summary': {'deep': {'count': 5, 'minutes': 6...		True
5	{'summary': {'deep': {'count': 5, 'minutes': 6...		True
6	{'summary': {'deep': None, 'wake': None, 'ligh...		False
7	{'summary': {'deep': {'count': 3, 'minutes': 6...		True
8	{'summary': {'deep': {'count': 3, 'minutes': 8...		True
9	{'summary': {'deep': {'count': 3, 'minutes': 8...		True
10	{'summary': {'deep': {'count': 2, 'minutes': 6...		True
11	{'summary': {'deep': {'count': 2, 'minutes': 6...		True
12	{'summary': {'deep': {'count': 4, 'minutes': 7...		True
13	{'summary': {'deep': {'count': 4, 'minutes': 7...		True
14	{'summary': {'deep': {'count': 7, 'minutes': 9...		True
15	{'summary': {'deep': {'count': 7, 'minutes': 9...		True
16	{'summary': {'deep': {'count': 7, 'minutes': 8...		True
17	{'summary': {'deep': {'count': 7, 'minutes': 8...		True
18	{'summary': {'deep': {'count': 3, 'minutes': 7...		True
19	{'summary': {'deep': {'count': 3, 'minutes': 7...		True
20	{'summary': {'deep': {'count': 3, 'minutes': 1...		True
21	{'summary': {'deep': {'count': 3, 'minutes': 1...		True
22	{'summary': {'deep': {'count': 7, 'minutes': 1...		True
23	{'summary': {'deep': {'count': 7, 'minutes': 1...		True
24	{'summary': {'deep': {'count': 6, 'minutes': 9...		True
25	{'summary': {'deep': {'count': 6, 'minutes': 9...		True
26	{'summary': {'deep': {'count': 5, 'minutes': 1...		True
27	{'summary': {'deep': {'count': 5, 'minutes': 1...		True
28	{'summary': {'deep': {'count': 3, 'minutes': 8...		True

```

29 {'summary': {'deep': {'count': 3, 'minutes': 8...      True
30 {'summary': {'deep': {'count': 6, 'minutes': 8...      True
31 {'summary': {'deep': {'count': 6, 'minutes': 8...      True

```

	logId	count	minutes	thirtyDayAvgMinutes
0	42724299942	5.0	104.0	0.0
1	42724299942	5.0	104.0	81.0
2	43071124476	5.0	87.0	80.0
3	43071124476	5.0	87.0	0.0
4	43415888071	5.0	67.0	88.0
5	43415888071	5.0	67.0	0.0
6	43483210808	NaN	NaN	NaN
7	43472303148	3.0	68.0	74.0
8	43758999257	3.0	80.0	77.0
9	43758999257	3.0	80.0	0.0
10	44096210623	2.0	62.0	0.0
11	44096210623	2.0	62.0	91.0
12	44452696017	4.0	79.0	79.0
13	44452696017	4.0	79.0	0.0
14	44796156145	7.0	92.0	0.0
15	44796156145	7.0	92.0	77.0
16	45141553280	7.0	85.0	74.0
17	45141553280	7.0	85.0	0.0
18	45472556013	3.0	75.0	80.0
19	45472556013	3.0	75.0	0.0
20	45804854753	3.0	100.0	73.0
21	45804854753	3.0	100.0	0.0
22	46126383325	7.0	107.0	73.0
23	46126383325	7.0	107.0	0.0
24	46439171327	6.0	90.0	0.0
25	46439171327	6.0	90.0	73.0
26	46747290318	5.0	110.0	81.0
27	46747290318	5.0	110.0	0.0
28	47046675063	3.0	85.0	0.0
29	47046675063	3.0	85.0	79.0
30	47353917202	6.0	84.0	0.0
31	47353917202	6.0	84.0	81.0

Very interesting, it looks like we have duplicates in the ‘levels’, specifically in the summary. The **thirtyDayAverageMinutes** seems to be miscomputed for some reason, we can handle this by flattening out the data structure. In addition it looks like we have just 1 record where the **mainSleep** was set to false (2023-11-14), but at least those records have different **logId** fields.

One very helpful function I just ran shows how awesome **DuckDB** is, the *unnest* function can break down json data into it’s individual parts. Depending on how the data is structured you can even refer to a lower level of data. In this case I was able to step into the **levels** field of data, into the **summary** subsection of **levels** and then grab one of the types of summaries for a sleep stage which was **deep** sleep and then flatten the data out. This is a very helpful and useful function when working with data and trying to reach important data nested in hierarchical structures.

What I'm going to do next is break up the data into different dataframes using **logId** as the key to link them all together, like setting up database tables with a link between them. I'll start with the outermost level of data, and for what I want to use this for, I don't need to look at sleep data that wasn't part of my main sleep.

```
[7]: sleep_meta_df = duckdb.sql("""Select distinct logId, dateOfSleep, startTime, \
    endTime, duration, minutesToFallAsleep, minutesAsleep, \
    minutesAwake, minutesAfterWakeup, timeInBed, \
    efficiency, type, infoCode, logType, mainSleep \
    FROM sleep_data \
    WHERE mainSleep == true""").df()
display(sleep_meta_df.head(6))
```

	logId	dateOfSleep	startTime	endTime	duration	\
0	45306958018	2024-04-22	2024-04-21 22:19:30	2024-04-22 05:31:30	25920000	
1	46987212286	2024-09-28	2024-09-27 22:12:00	2024-09-28 07:05:30	31980000	
2	46876496336	2024-09-17	2024-09-16 23:24:00	2024-09-17 06:12:00	24480000	
3	46843361089	2024-09-14	2024-09-13 22:03:30	2024-09-14 04:32:30	23340000	
4	46806686902	2024-09-10	2024-09-09 22:13:30	2024-09-10 06:04:30	28260000	
5	46795305204	2024-09-09	2024-09-08 21:32:00	2024-09-09 06:09:30	31020000	

	minutesToFallAsleep	minutesAsleep	minutesAwake	minutesAfterWakeup	\
0	0	388	44	2	
1	0	480	53	1	
2	0	358	50	0	
3	0	354	35	3	
4	0	427	44	0	
5	0	464	53	1	

	timeInBed	efficiency	type	infoCode	logType	mainSleep
0	432	98	stages	0	auto_detected	True
1	533	100	stages	0	auto_detected	True
2	408	99	stages	0	auto_detected	True
3	389	98	stages	0	auto_detected	True
4	471	99	stages	0	auto_detected	True
5	517	99	stages	0	auto_detected	True

By using the *distinct* clause on the outermost level of data I can eliminate all of the duplicates except for the one scenario where I had 2 records for the same night, for that I just filtered the data in the *where* clause for **mainSleep** == true.

Next I want the sleep summary data, the problem is I will need to first unnest the data from the levels.summary, but each stage of sleep (deep, rem, light, & awake) has it's own set of data and the key is the name of the stage. I want the name of the stage as a data value, as well as it's dataset, so what I've done here is build a couple of CTEs. First I unnest the first level, then I break down the different layers and add the key in for each and call that variable "stage". All of the stages have the same data (count, minutes, thirtyDayAverageMinutes) which is why this works so easily. Finally in order to get rid of the data problems we saw earlier where the thirtyDayAverages were zeros I just take the max for those, which avoids the duplicates and gives me my data.

```
[8]: sleep_summary = duckdb.sql("""with summary_data as (select logId, unnest(levels.
↳summary, max_depth := 1) from sleep_data),
    unioned_data as (
        select logId, 'deep' as stage, unnest(deep) from summary_data
        union all
        select logId, 'wake' as stage, unnest(wake) from summary_data
        union all
        select logId, 'light' as stage, unnest(light) from summary_data
        union all
        select logId, 'rem' as stage, unnest(rem) from summary_data)
    select logId, stage, count, minutes, max(thirtyDayAvgMinutes) as
↳thirtyDayAvgMinutes
    from unioned_data
    group by logId, stage, count, minutes
    order by logId, stage""")
sleep_summary_df = sleep_summary.df()
display(sleep_summary_df.head(12))
```

	logId	stage	count	minutes	thirtyDayAvgMinutes
0	42400922920	deep	3.0	66.0	0.0
1	42400922920	light	30.0	272.0	0.0
2	42400922920	rem	7.0	78.0	0.0
3	42400922920	wake	28.0	42.0	0.0
4	42412586924	deep	4.0	86.0	66.0
5	42412586924	light	29.0	218.0	272.0
6	42412586924	rem	7.0	87.0	78.0
7	42412586924	wake	26.0	34.0	42.0
8	42424336860	deep	2.0	58.0	76.0
9	42424336860	light	30.0	260.0	245.0
10	42424336860	rem	7.0	93.0	83.0
11	42424336860	wake	31.0	54.0	38.0

Alright now we're getting somewhere! I could stop here as this is probably the level of detail I will want for my predictions, but let's keep going and get the other data. The next set of data is what makes up the summary, theoretically if we total up the data it should equal what we see above. Let's get the data first and then aggregate it.

```
[9]: sleep_data_detail = duckdb.sql("""select logId, unnest(levels.data, recursive :
↳= True) from sleep_data order by logId""")
display(sleep_data_detail.df().head(12))
```

	logId	dateTime	level	seconds
0	42400922920	2023-08-12 22:16:00	wake	330
1	42400922920	2023-08-12 22:21:30	light	3960
2	42400922920	2023-08-12 23:27:30	rem	630
3	42400922920	2023-08-12 23:38:00	light	720
4	42400922920	2023-08-12 23:50:00	deep	3180
5	42400922920	2023-08-13 00:43:00	light	540
6	42400922920	2023-08-13 00:52:00	rem	930

7	42400922920	2023-08-13 01:07:30	light	3030
8	42400922920	2023-08-13 01:58:00	wake	360
9	42400922920	2023-08-13 02:04:00	light	180
10	42400922920	2023-08-13 02:07:00	rem	1560
11	42400922920	2023-08-13 02:33:00	light	3390

```
[10]: sleep_data_agg = duckdb.sql("""select logId, level, count(*) as
↳count_of_stages, sum(seconds)/60 as total_minutes from sleep_data_detail
      group by logId, level
      order by logId, level""")
display(sleep_data_agg.df().head(12))
```

	logId	level	count_of_stages	total_minutes
0	42400922920	deep	3	66.5
1	42400922920	light	10	293.5
2	42400922920	rem	5	79.5
3	42400922920	wake	3	19.0
4	42412586924	deep	4	86.0
5	42412586924	light	10	240.0
6	42412586924	rem	6	88.0
7	42412586924	wake	2	11.5
8	42424336860	deep	2	59.0
9	42424336860	light	11	279.5
10	42424336860	rem	5	97.5
11	42424336860	wake	5	29.0

I definitely did not expect this but it looks like the data in the summary doesn't match the data supplied in the **data** field if you aggregate it. I can understand the time not matching up as the data looks like it only logs 30 second intervals, but the counts are off which is weird. Maybe the **shortData** field is also included and then both the **shortData** and **data** fields are aggregated, let's try that.

```
[11]: sleep_short_data = duckdb.sql("""select logId, unnest(levels.shortData,
↳recursive := True) from sleep_data order by logId""")
display(sleep_short_data.df().head(10))
sleep_data_all_agg = duckdb.sql("""WITH full_sleep_data AS (
      select * from sleep_data_detail
      union all
      select * from sleep_short_data)
      select logId, level as stage, count(*) as
↳count_of_stages, sum(seconds)/60 as total_minutes from full_sleep_data
      group by logId, level
      order by logId, level""")
display(sleep_data_all_agg.df().head(12))
```

	logId	dateTime	level	seconds
0	42400922920	2023-08-12 22:28:00	wake	30
1	42400922920	2023-08-12 23:11:00	wake	30
2	42400922920	2023-08-12 23:23:00	wake	30

3	42400922920	2023-08-12	23:40:30	wake	90
4	42400922920	2023-08-12	23:44:00	wake	30
5	42400922920	2023-08-13	00:42:00	wake	60
6	42400922920	2023-08-13	00:50:30	wake	30
7	42400922920	2023-08-13	01:04:00	wake	30
8	42400922920	2023-08-13	01:10:00	wake	30
9	42400922920	2023-08-13	01:16:30	wake	60

	logId	stage	count_of_stages	total_minutes
0	42400922920	deep	3	66.5
1	42400922920	light	10	293.5
2	42400922920	rem	5	79.5
3	42400922920	wake	28	42.5
4	42412586924	deep	4	86.0
5	42412586924	light	10	240.0
6	42412586924	rem	6	88.0
7	42412586924	wake	26	34.5
8	42424336860	deep	2	59.0
9	42424336860	light	11	279.5
10	42424336860	rem	5	97.5
11	42424336860	wake	31	54.0

It looks like that fixed the counts with **wake** but not much else, for a decent understanding of how different the data is we can join the 2 sets together and then take the differences. We can plot the data just to get an idea of how different the data is for both count and minutes. Like I mentioned before I am going to stick with the summary data as I have a feeling it's probably more accurate as it is precomputed. I can see the actual data is only accurate to a 30 second interval.

```
[12]: sleep_data_diff = duckdb.sql("""select s.logId, s.stage, s.count as
    ↳summary_count_of_stages, s.minutes as summary_total_minutes,
    ↳d.count_of_stages as detail_count_of_stages, d.
    ↳total_minutes as detail_total_minutes,
    ↳s.count - d.count_of_stages as
    ↳summary_less_detail_count, s.minutes - d.total_minutes as
    ↳summary_less_detail_minutes
    from sleep_summary s
    inner join sleep_data_all_agg d
    on s.logId = d.logId
    and s.stage = d.stage
    order by s.logId, s.stage""").df()
display(sleep_data_diff.head(12))
```

	logId	stage	summary_count_of_stages	summary_total_minutes	\
0	42400922920	deep	3	66	
1	42400922920	light	30	272	
2	42400922920	rem	7	78	
3	42400922920	wake	28	42	
4	42412586924	deep	4	86	
5	42412586924	light	29	218	

6	42412586924	rem	7	87
7	42412586924	wake	26	34
8	42424336860	deep	2	58
9	42424336860	light	30	260
10	42424336860	rem	7	93
11	42424336860	wake	31	54

	detail_count_of_stages	detail_total_minutes	summary_less_detail_count	\
0	3	66.5	0	
1	10	293.5	20	
2	5	79.5	2	
3	28	42.5	0	
4	4	86.0	0	
5	10	240.0	19	
6	6	88.0	1	
7	26	34.5	0	
8	2	59.0	0	
9	11	279.5	19	
10	5	97.5	2	
11	31	54.0	0	

	summary_less_detail_minutes
0	-0.5
1	-21.5
2	-1.5
3	-0.5
4	0.0
5	-22.0
6	-1.0
7	-0.5
8	-1.0
9	-19.5
10	-4.5
11	0.0


```
[13]: import seaborn as sns
import matplotlib.pyplot as plt
sns.set_style("ticks")
sns.set_context("notebook")
g = sns.catplot(data=sleep_data_diff,x="stage", y="summary_less_detail_count",
    kind="box", height=6, aspect=1.5)
g.set_axis_labels("Sleep Stages", "Difference (Summary - Detail)")
g.fig.suptitle('Difference in Sleep Stage Counts')
g.fig.subplots_adjust(top=.93)
plt.show()
g = sns.catplot(data=sleep_data_diff,x="stage",
    y="summary_less_detail_minutes", kind="box", height=6, aspect=1.5)
```

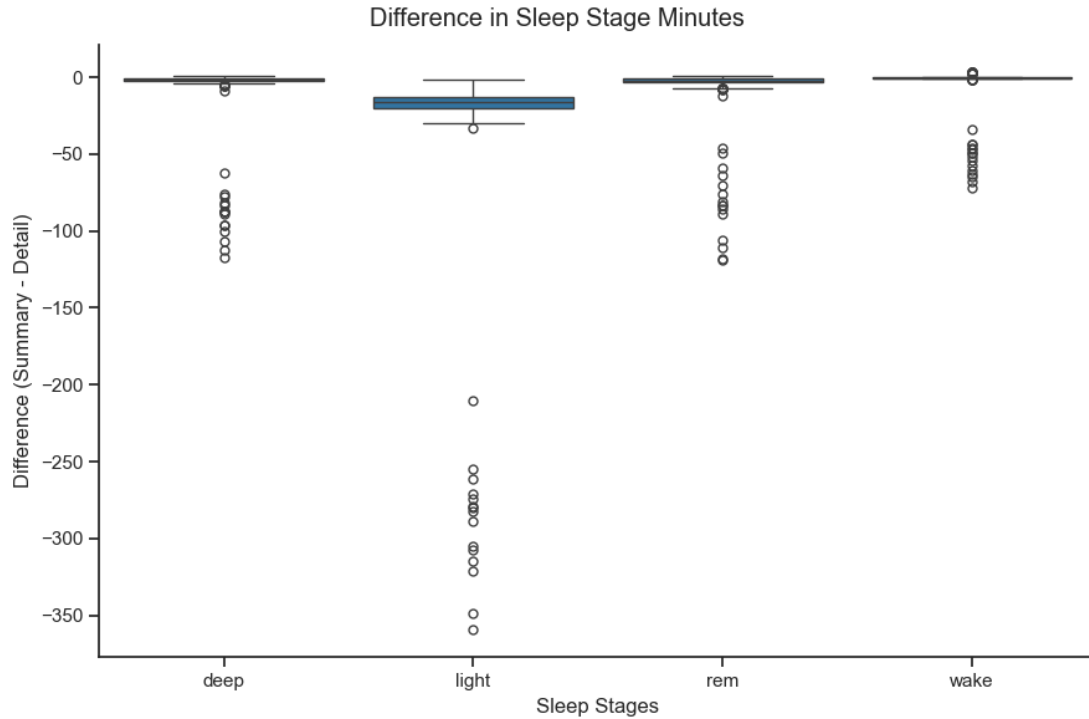


```

g.set_axis_labels("Sleep Stages", "Difference (Summary - Detail)")
g.fig.suptitle('Difference in Sleep Stage Minutes')
g.fig.subplots_adjust(top=.93)
plt.show()

```





It's clear with these 2 plots that our data doesn't really agree for *light* sleep stages, *deep* & *wake*, seem tolerable and *rem* is okay. That's all we need to know, I figured I'd break up the monotony of text, code, & data with a plot or 2. Now onto other datasets.

1.3 Heart Rate & HRV

We are going to be moving a little faster here so what we want to do is bring in the heart rate data. Both the actual measures of heart rate and then the heart rate variability data. The **heart rate data** is in a json format, which has a **bpm** and **confidence** measure for each timestamp. The **confidence** measure goes from 0 (no heart rate detected) to 3 (high confidence). The **dateTime** intervals seem to be irregular measures between 5 and 10 seconds. The **HRV Data** is a CSV and is a summary set of data that is gathered once a day, it represents the numbers you see on your sleep metrics for the Heart Rate Variability. There are other more frequent measures of HRV but it's not clear how they get this summary number from those so I'm sticking with the summary. The **rmssd** is the field that we are after for HRV and it's important to note that it takes this measure during sleep and will report it essentially after you wake up, but it gets assigned to that date too. For example if I start sleep on 2024-12-22 at 22:00:00 and wake up on 2024-12-23 at 06:00:00, my HRV score will be measured during sleep and assigned with the 2024-12-23 date.

1.4 Establishing The Right Dataset for Predictions

It is extremely important to understand if my goal is to create a dataset to predict HRV, I want to make this as realistic as possible so using my example date above, I should be taking data from the day of 2024-12-22 to predict the HRV I was assigned on 2024-12-23. Also if I truly wanted to make this a good quality model, I would also want to make sure I do not take any data while I am

sleeping, that's kind of like cheating. I want to be able to predict something BEFORE it happens not during the event where it's being measured. In order to make this a strong model I am going to filter my data to 2 hours before I fall asleep, this way someone could actually go get FitBit data from Google, download it and insert it into the model to make a prediction for the HRV score. This also means a lot of data wrangling and filtering which is why I chose DuckDB as I know how to manipulate data a lot easier with SQL than I do with Pandas.

```
[14]: heart_rate_data = duckdb.read_json('./data/unzipped/Takeout/Fitbit/Global_
↳Export Data/heart_rate-*.json', timestamp_format="%m/%d/%y %H:%M:%S")
heart_rate_df = duckdb.sql('select dateTime, unnest(value) from_
↳heart_rate_data').df()
display(heart_rate_df.head(10))

hrv_df = duckdb.read_csv('./data/unzipped/Takeout/Fitbit/Heart Rate Variability/
↳Daily Heart Rate Variability Summary*.csv', timestamp_format="%Y-%m-%dT%H:%M:
↳%S").df()
display(hrv_df.sort_values('timestamp').head(10))
```

	dateTime	bpm	confidence
0	2023-08-12 18:36:31	70	0
1	2023-08-12 18:36:46	64	0
2	2023-08-12 18:36:51	63	1
3	2023-08-12 18:36:56	65	1
4	2023-08-12 18:37:01	72	1
5	2023-08-12 18:37:06	76	1
6	2023-08-12 18:37:16	74	2
7	2023-08-12 18:37:21	73	2
8	2023-08-12 18:37:31	71	1
9	2023-08-12 18:37:36	70	1

	timestamp	rmssd	nremhr	entropy
0	2023-08-13	48.934	60.798	2.855
125	2023-08-14	40.485	67.520	2.632
364	2023-08-15	51.453	58.189	2.803
17	2023-08-16	37.462	64.676	2.446
336	2023-08-17	53.514	60.341	2.801
304	2023-08-18	48.468	62.162	2.742
394	2023-08-19	52.700	56.938	2.955
14	2023-08-20	54.221	61.383	3.103
2	2023-08-21	65.983	54.468	2.953
3	2023-08-22	50.733	57.393	2.787

In order to prepare the sleep summary data I am going to have to pivot it from a long format (multiple rows per observation) to a wide format (1 row per observation). There are a number of ways to do this in both DuckDB and Pandas. I am going to show you how to produce the same dataset by using *pivot* functions and then time the performance of both. Note that pandas is faster but in my opinion DuckDB is more readable. I only have the data that I have which is ~1800 rows. Perhaps if there was a lot more data DuckDB would be faster...perhaps not, all I can tell you is that not everything will be faster with DuckDB.

```
[15]: %%timeit
adj_sleep_summary_pandas_df =
    ↪sleep_summary_df[['logId', 'stage', 'count', 'minutes', 'thirtyDayAvgMinutes']].
    ↪pivot(index='logId', columns='stage')
adj_sleep_summary_pandas_df.columns = [f"{col[1]}_{col[0]}" for col in
    ↪adj_sleep_summary_pandas_df.columns]

# Reset index for a regular looking DataFrame
adj_sleep_summary_pandas_df.reset_index(inplace=True)
```

1.32 ms ± 45 s per loop (mean ± std. dev. of 7 runs, 1,000 loops each)

```
[16]: %%timeit
adj_sleep_summary_duck_df = duckdb.sql("""pivot sleep_summary on stage
    using
        first(count) as count,
        first(minutes) as minutes,
        first(thirtyDayAvgMinutes) as thirtyDayAvgMinutes
    order by logId""")
```

423 ms ± 22.6 ms per loop (mean ± std. dev. of 7 runs, 1 loop each)

```
[17]: adj_sleep_summary_df = duckdb.sql("""pivot sleep_summary on stage
    using
        first(count) as count,
        first(minutes) as minutes,
        first(thirtyDayAvgMinutes) as thirtyDayAvgMinutes
    order by logId""").df()
display(adj_sleep_summary_df.head(10))
```

	logId	deep_count	deep_minutes	deep_thirtyDayAvgMinutes	\
0	42400922920	3.0	66.0	0.0	
1	42412586924	4.0	86.0	66.0	
2	42424336860	2.0	58.0	76.0	
3	42437088181	4.0	128.0	70.0	
4	42449293365	5.0	87.0	85.0	
5	42463261545	5.0	93.0	85.0	
6	42476649087	4.0	47.0	86.0	
7	42487823141	4.0	77.0	81.0	
8	42498649806	4.0	77.0	80.0	
9	42510611886	5.0	87.0	80.0	

	light_count	light_minutes	light_thirtyDayAvgMinutes	rem_count	\
0	30.0	272.0	0.0	7.0	
1	29.0	218.0	272.0	7.0	
2	30.0	260.0	245.0	7.0	
3	28.0	204.0	250.0	7.0	
4	23.0	228.0	239.0	9.0	
5	31.0	265.0	236.0	12.0	

6	34.0	284.0	241.0	10.0
7	30.0	276.0	247.0	7.0
8	27.0	242.0	251.0	6.0
9	27.0	222.0	250.0	6.0

	rem_minutes	rem_thirtyDayAvgMinutes	wake_count	wake_minutes	\
0	78.0	0.0	28.0	42.0	
1	87.0	78.0	26.0	34.0	
2	93.0	83.0	31.0	54.0	
3	72.0	86.0	24.0	51.0	
4	70.0	83.0	27.0	51.0	
5	81.0	80.0	32.0	34.0	
6	82.0	80.0	35.0	42.0	
7	75.0	80.0	29.0	47.0	
8	105.0	80.0	27.0	39.0	
9	86.0	83.0	28.0	42.0	

	wake_thirtyDayAvgMinutes
0	0.0
1	42.0
2	38.0
3	43.0
4	45.0
5	46.0
6	44.0
7	44.0
8	44.0
9	44.0

1.5

```
[18]: print(adj_sleep_summary_df.shape[0])
      print(sleep_meta_df.shape[0])
```

466
465

```
[19]: complete_sleep_data = sleep_meta_df.merge(adj_sleep_summary_df, on='logId').
      ↪sort_values('dateOfSleep')
      display(complete_sleep_data.tail(5))
```

	logId	dateOfSleep	startTime	endTime	\
71	47560780256	2024-11-23	2024-11-23 00:19:00	2024-11-23 07:32:30	
423	47568889247	2024-11-24	2024-11-23 22:22:00	2024-11-24 07:31:30	
376	47578085269	2024-11-25	2024-11-24 22:42:30	2024-11-25 05:45:30	
375	47589819717	2024-11-26	2024-11-25 22:29:30	2024-11-26 06:44:30	
70	47599022157	2024-11-27	2024-11-26 22:39:00	2024-11-27 04:01:30	

	duration	minutesToFallAsleep	minutesAsleep	minutesAwake	\
--	----------	---------------------	---------------	--------------	---

71	25980000	0	379	54
423	32940000	0	488	61
376	25380000	0	372	51
375	29700000	0	457	38
70	19320000	0	290	32

	minutesAfterWakeup	timeInBed	...	deep_thirtyDayAvgMinutes	light_count	\
71	0	433	...	71.0	22.0	
423	0	549	...	70.0	38.0	
376	0	423	...	71.0	19.0	
375	0	495	...	72.0	21.0	
70	0	322	...	72.0	14.0	

	light_minutes	light_thirtyDayAvgMinutes	rem_count	rem_minutes	\
71	279.0	270.0	7.0	48.0	
423	314.0	270.0	7.0	75.0	
376	233.0	272.0	5.0	64.0	
375	262.0	271.0	7.0	104.0	
70	202.0	270.0	2.0	39.0	

	rem_thirtyDayAvgMinutes	wake_count	wake_minutes	\
71	70.0	23.0	54.0	
423	69.0	35.0	61.0	
376	69.0	20.0	51.0	
375	69.0	23.0	38.0	
70	70.0	13.0	32.0	

	wake_thirtyDayAvgMinutes
71	52.0
423	52.0
376	52.0
375	52.0
70	51.0

[5 rows x 27 columns]

If you're wondering why I didn't just use the SQL syntax below...I was just being lazy, with pandas if you use a merge (effectively an inner join), if you have the same name for the **on** column, it won't be repeated in the dataset. I could either write the SQL join and type out all the column names except for the second **logID** column or I could just use pandas, so I used pandas.

```
[20]: complete_sleep_data = sleep_meta_df.merge(adj_sleep_summary_df, on='logId').
      ↪sort_values('dateOfSleep')
      display(complete_sleep_data.tail(5))
```

	logId	dateOfSleep	startTime	endTime	\
71	47560780256	2024-11-23	2024-11-23 00:19:00	2024-11-23 07:32:30	
423	47568889247	2024-11-24	2024-11-23 22:22:00	2024-11-24 07:31:30	

376	47578085269	2024-11-25	2024-11-24 22:42:30	2024-11-25 05:45:30
375	47589819717	2024-11-26	2024-11-25 22:29:30	2024-11-26 06:44:30
70	47599022157	2024-11-27	2024-11-26 22:39:00	2024-11-27 04:01:30

	duration	minutesToFallAsleep	minutesAsleep	minutesAwake	\
71	25980000	0	379	54	
423	32940000	0	488	61	
376	25380000	0	372	51	
375	29700000	0	457	38	
70	19320000	0	290	32	

	minutesAfterWakeup	timeInBed	...	deep_thirtyDayAvgMinutes	light_count	\
71	0	433	...	71.0	22.0	
423	0	549	...	70.0	38.0	
376	0	423	...	71.0	19.0	
375	0	495	...	72.0	21.0	
70	0	322	...	72.0	14.0	

	light_minutes	light_thirtyDayAvgMinutes	rem_count	rem_minutes	\
71	279.0	270.0	7.0	48.0	
423	314.0	270.0	7.0	75.0	
376	233.0	272.0	5.0	64.0	
375	262.0	271.0	7.0	104.0	
70	202.0	270.0	2.0	39.0	

	rem_thirtyDayAvgMinutes	wake_count	wake_minutes	\
71	70.0	23.0	54.0	
423	69.0	35.0	61.0	
376	69.0	20.0	51.0	
375	69.0	23.0	38.0	
70	70.0	13.0	32.0	

	wake_thirtyDayAvgMinutes
71	52.0
423	52.0
376	52.0
375	52.0
70	51.0

[5 rows x 27 columns]

1.6 Final Aggregation

Here is the final manipulation of the data where I am trying to establish the right dataset for a practical predictive model. I am using CTEs, the first is based on the sleep data which is essentially 1 record per day. I use the sleep data to establish the proper dates and times. I am using the **startTime** less 2 hours to the next row's **startTime** less 2 hours. From a temporal perspective the data starts 2 hours before I sleep, continues through my sleep and ends 2 hours

before I fall asleep the next day. The **dateOfSleep** field records the date you woke up, so I want to predict the NEXT days HRV which is measured and computed during that sleeping interval. All of the heart rate data I want to use and aggregate is set by those start times which you can see in the second CTE, which is basically a cross join but where the heart rate is between the 2 start times established by the sleep data. The third and final CTE is used to aggregate that data so we have 1 record per sleep, so what we do is take measures using *max*, *min*, *std dev*, *median*, *10th percentile*, *90th percentile*, and the count of my bpm data, along with all the data from the sleeping records. Also just to see from a temporal perspective how well the previous HRV predicts the next days HRV I include that as a feature in the final query.

```
[21]: final_df = duckdb.sql("""WITH Complete_Sleep AS (
    select *,
        (dateOfSleep::Date + 1) as hrv_prediction_date,
        (startTime - Interval '2 Hours') as
↪startTime_sleep_less_2_hrs,
        ((lead(startTime) over (order by logId)) - Interval '2
↪Hours') as tomorrow_startTime_sleep_less_2_hrs
    from complete_sleep_data),
    HR_with_sleep AS (
        Select hr.*, cs.* from Complete_Sleep cs, heart_rate_df hr
        where hr.dateTime >= cs.startTime_sleep_less_2_hrs
        and hr.dateTime < tomorrow_startTime_sleep_less_2_hrs),
    hr_sleep_agg AS (
        Select dateOfSleep::date as dateOfSleep, hrv_prediction_date,
        max(bpm) as bpm_max, min(bpm) as bpm_min, round(avg(bpm), 2) as
↪bpm_avg, round(stddev_samp(bpm), 2) as bpm_std,
        round(median(bpm),2) as bpm_median, round(quantile_cont(bpm, 0.
↪10),2) as bpm_10th, round(quantile_cont(bpm, 0.90),2) as bpm_90th,
        count(bpm) as bpm_count, minutesAsleep,
        minutesAwake,
        timeInBed,
        deep_count,
        light_count,
        rem_count,
        wake_count,
        deep_minutes,
        light_minutes,
        rem_minutes,
        wake_minutes,
        deep_thirtyDayAvgMinutes,
        light_thirtyDayAvgMinutes,
        rem_thirtyDayAvgMinutes,
        wake_thirtyDayAvgMinutes
    from HR_with_sleep
    Group by dateOfSleep::date, hrv_prediction_date, minutesAsleep,
↪minutesAwake,
        timeInBed,
```



```

        deep_count,
        light_count,
        rem_count,
        wake_count,
        deep_minutes,
        light_minutes,
        rem_minutes,
        wake_minutes,
        deep_thirtyDayAvgMinutes,
        light_thirtyDayAvgMinutes,
        rem_thirtyDayAvgMinutes,
        wake_thirtyDayAvgMinutes
    )

    select agg.*, lag(rmssd) over (order by dateOfSleep) as prev_hrv, hrv.
    ↳rmssd as target_hrv from hr_sleep_agg agg inner join hrv_df hrv on agg.
    ↳hrv_prediction_date = hrv.timestamp Order by dateOfSleep
    """).df()

```

```

[22]: pd.set_option('display.max_columns', None)
      final_df.drop("hrv_prediction_date", axis=1, inplace=True)
      display(final_df)

```

	dateOfSleep	bpm_max	bpm_min	bpm_avg	bpm_std	bpm_median	bpm_10th	\
0	2023-08-13	141	49	75.77	14.87	75.0	59.0	
1	2023-08-14	121	50	80.37	12.20	82.0	64.0	
2	2023-08-15	168	49	79.51	21.67	77.0	57.0	
3	2023-08-16	140	49	82.52	16.05	84.0	61.0	
4	2023-08-17	119	51	73.48	11.42	75.0	58.0	
..	
449	2024-11-21	122	48	72.00	13.56	73.0	55.0	
450	2024-11-22	161	53	82.70	21.17	79.0	59.0	
451	2024-11-23	105	51	71.38	11.18	71.0	58.0	
452	2024-11-24	165	46	78.99	29.02	68.0	52.0	
453	2024-11-25	130	47	71.56	14.60	70.0	54.0	

	bpm_90th	bpm_count	minutesAsleep	minutesAwake	timeInBed	deep_count	\
0	96.0	11070	416	42	458	3.0	
1	95.0	10169	391	34	425	4.0	
2	105.0	10797	411	54	465	2.0	
3	103.0	10225	404	51	455	4.0	
4	87.0	10522	385	51	436	5.0	
..	
449	88.0	10101	379	50	429	3.0	
450	115.0	11472	364	56	420	4.0	
451	86.0	9222	379	54	433	3.0	
452	129.0	12668	488	61	549	5.0	
453	92.0	9869	372	51	423	3.0	

	light_count	rem_count	wake_count	deep_minutes	light_minutes	\
0	30.0	7.0	28.0	66.0	272.0	
1	29.0	7.0	26.0	86.0	218.0	
2	30.0	7.0	31.0	58.0	260.0	
3	28.0	7.0	24.0	128.0	204.0	
4	23.0	9.0	27.0	87.0	228.0	
..	
449	20.0	5.0	19.0	79.0	248.0	
450	23.0	3.0	20.0	62.0	237.0	
451	22.0	7.0	23.0	52.0	279.0	
452	38.0	7.0	35.0	99.0	314.0	
453	19.0	5.0	20.0	75.0	233.0	

	rem_minutes	wake_minutes	deep_thirtyDayAvgMinutes	\
0	78.0	42.0	0.0	
1	87.0	34.0	66.0	
2	93.0	54.0	76.0	
3	72.0	51.0	70.0	
4	70.0	51.0	85.0	
..	
449	52.0	50.0	71.0	
450	65.0	56.0	72.0	
451	48.0	54.0	71.0	
452	75.0	61.0	70.0	
453	64.0	51.0	71.0	

	light_thirtyDayAvgMinutes	rem_thirtyDayAvgMinutes	\
0	0.0	0.0	
1	272.0	78.0	
2	245.0	83.0	
3	250.0	86.0	
4	239.0	83.0	
..	
449	273.0	71.0	
450	272.0	70.0	
451	270.0	70.0	
452	270.0	69.0	
453	272.0	69.0	

	wake_thirtyDayAvgMinutes	prev_hrv	target_hrv
0	0.0	NaN	40.485
1	42.0	40.485	51.453
2	38.0	51.453	37.462
3	43.0	37.462	53.514
4	45.0	53.514	48.468
..
449	51.0	51.494	42.113
450	51.0	42.113	44.060

451	52.0	44.060	58.188
452	52.0	58.188	60.853
453	52.0	60.853	49.758

[454 rows x 26 columns]

1.7 Bonus: Correlation Filtering

When you're preparing data for a model you generally want to exclude data that is highly correlated. "Highly" correlated is ambiguous so I'm going to use 0.9 as a threshold. You can do this with code automatically but I chose the manual approach to show you. I also show both the correlation matrix and a plot, it's much easier to just see the data and find the correlations than it is to read a square matrix of numbers.

```
[23]: final_df.corr()
```

```
[23]:
```

	dateOfSleep	bpm_max	bpm_min	bpm_avg	\
dateOfSleep	1.000000	-0.034703	-0.222612	-0.233377	
bpm_max	-0.034703	1.000000	-0.032256	0.616479	
bpm_min	-0.222612	-0.032256	1.000000	0.435690	
bpm_avg	-0.233377	0.616479	0.435690	1.000000	
bpm_std	0.039252	0.862258	-0.222124	0.601060	
bpm_median	-0.342522	0.174308	0.495046	0.785332	
bpm_10th	-0.300051	0.002621	0.852287	0.551725	
bpm_90th	0.006009	0.770030	0.036941	0.789406	
bpm_count	-0.012505	0.392694	-0.014583	0.358291	
minutesAsleep	0.019802	0.019281	-0.208599	-0.230068	
minutesAwake	0.038499	-0.111518	-0.134464	-0.129557	
timeInBed	0.026381	-0.007233	-0.223246	-0.244961	
deep_count	-0.113175	-0.039884	-0.076466	-0.105497	
light_count	-0.071701	0.037702	-0.167459	-0.130345	
rem_count	-0.182631	0.017710	0.161166	0.057375	
wake_count	-0.136110	0.069070	-0.096298	-0.064391	
deep_minutes	-0.116257	0.040020	-0.019692	0.029081	
light_minutes	0.128219	-0.014968	-0.317837	-0.262923	
rem_minutes	-0.052281	0.066116	0.059835	-0.002994	
wake_minutes	0.044509	-0.101243	-0.149808	-0.101802	
deep_thirtyDayAvgMinutes	-0.092159	0.011244	0.110658	0.155230	
light_thirtyDayAvgMinutes	0.222209	-0.048390	-0.204543	-0.219024	
rem_thirtyDayAvgMinutes	-0.046121	-0.021180	0.031532	0.031240	
wake_thirtyDayAvgMinutes	0.186204	-0.026494	-0.006736	-0.053862	
prev_hrv	0.025463	0.027929	-0.729785	-0.402024	
target_hrv	0.025688	-0.294057	-0.225867	-0.318599	

	bpm_std	bpm_median	bpm_10th	bpm_90th	\
dateOfSleep	0.039252	-0.342522	-0.300051	0.006009	
bpm_max	0.862258	0.174308	0.002621	0.770030	
bpm_min	-0.222124	0.495046	0.852287	0.036941	

bpm_avg	0.601060	0.785332	0.551725	0.789406
bpm_std	1.000000	0.108387	-0.216658	0.911308
bpm_median	0.108387	1.000000	0.603565	0.315115
bpm_10th	-0.216658	0.603565	1.000000	0.077574
bpm_90th	0.911308	0.315115	0.077574	1.000000
bpm_count	0.460249	0.070436	0.025891	0.500103
minutesAsleep	0.081034	-0.362392	-0.215607	-0.010050
minutesAwake	-0.017095	-0.144526	-0.126706	-0.039130
timeInBed	0.072211	-0.374562	-0.229250	-0.017987
deep_count	-0.017979	-0.108674	-0.073619	-0.049728
light_count	0.111178	-0.243340	-0.177767	0.044032
rem_count	-0.025332	0.002315	0.208387	-0.005131
wake_count	0.109687	-0.175163	-0.079483	0.050090
deep_minutes	0.042664	-0.025979	0.065534	0.041224
light_minutes	0.106530	-0.369781	-0.370895	0.007787
rem_minutes	0.034241	-0.082595	0.068026	0.015588
wake_minutes	0.008633	-0.124105	-0.137978	-0.009024
deep_thirtyDayAvgMinutes	-0.009105	0.186094	0.185833	0.036697
light_thirtyDayAvgMinutes	0.015408	-0.243644	-0.296932	-0.057919
rem_thirtyDayAvgMinutes	-0.009556	0.029352	0.056902	-0.002272
wake_thirtyDayAvgMinutes	-0.036509	-0.037234	-0.028147	-0.051236
prev_hrv	0.181155	-0.393489	-0.810916	-0.073119
target_hrv	-0.208096	-0.217225	-0.237110	-0.235557

	bpm_count	minutesAsleep	minutesAwake	timeInBed \
dateOfSleep	-0.012505	0.019802	0.038499	0.026381
bpm_max	0.392694	0.019281	-0.111518	-0.007233
bpm_min	-0.014583	-0.208599	-0.134464	-0.223246
bpm_avg	0.358291	-0.230068	-0.129557	-0.244961
bpm_std	0.460249	0.081034	-0.017095	0.072211
bpm_median	0.070436	-0.362392	-0.144526	-0.374562
bpm_10th	0.025891	-0.215607	-0.126706	-0.229250
bpm_90th	0.500103	-0.010050	-0.039130	-0.017987
bpm_count	1.000000	0.381207	0.169575	0.396812
minutesAsleep	0.381207	1.000000	0.161908	0.976578
minutesAwake	0.169575	0.161908	1.000000	0.364057
timeInBed	0.396812	0.976578	0.364057	1.000000
deep_count	0.104383	0.334886	0.094187	0.333569
light_count	0.233793	0.605575	0.260005	0.629663
rem_count	0.115941	0.320341	-0.132725	0.276157
wake_count	0.215900	0.583821	0.166249	0.589330
deep_minutes	0.237387	0.496881	0.025971	0.477823
light_minutes	0.200950	0.682627	0.315229	0.715528
rem_minutes	0.180984	0.509290	-0.299799	0.419898
wake_minutes	0.135607	0.118121	1.000000	0.328420
deep_thirtyDayAvgMinutes	-0.002358	-0.006983	0.028399	-0.002073
light_thirtyDayAvgMinutes	-0.020538	-0.025883	0.012103	-0.017975

rem_thirtyDayAvgMinutes	-0.073252	0.000271	-0.024911	0.002071
wake_thirtyDayAvgMinutes	0.011239	-0.037440	0.054628	-0.028304
prev_hrv	0.005807	0.116860	0.089138	0.128224
target_hrv	-0.264832	-0.101370	0.052049	-0.083061

	deep_count	light_count	rem_count	wake_count	\
dateOfSleep	-0.113175	-0.071701	-0.182631	-0.136110	
bpm_max	-0.039884	0.037702	0.017710	0.069070	
bpm_min	-0.076466	-0.167459	0.161166	-0.096298	
bpm_avg	-0.105497	-0.130345	0.057375	-0.064391	
bpm_std	-0.017979	0.111178	-0.025332	0.109687	
bpm_median	-0.108674	-0.243340	0.002315	-0.175163	
bpm_10th	-0.073619	-0.177767	0.208387	-0.079483	
bpm_90th	-0.049728	0.044032	-0.005131	0.050090	
bpm_count	0.104383	0.233793	0.115941	0.215900	
minutesAsleep	0.334886	0.605575	0.320341	0.583821	
minutesAwake	0.094187	0.260005	-0.132725	0.166249	
timeInBed	0.333569	0.629663	0.276157	0.589330	
deep_count	1.000000	0.342781	0.102288	0.194909	
light_count	0.342781	1.000000	0.174682	0.890613	
rem_count	0.102288	0.174682	1.000000	0.389770	
wake_count	0.194909	0.890613	0.389770	1.000000	
deep_minutes	0.359378	0.250720	0.292509	0.294964	
light_minutes	0.125919	0.579264	-0.089013	0.486325	
rem_minutes	0.169642	0.077876	0.555045	0.147443	
wake_minutes	0.094187	0.260005	-0.132725	0.166249	
deep_thirtyDayAvgMinutes	0.077491	0.053393	0.057468	0.092272	
light_thirtyDayAvgMinutes	-0.052548	-0.003848	-0.099303	-0.025634	
rem_thirtyDayAvgMinutes	0.000782	-0.024750	0.035218	0.023740	
wake_thirtyDayAvgMinutes	0.040581	-0.056804	-0.101375	-0.070096	
prev_hrv	0.128256	0.139161	-0.128832	0.083330	
target_hrv	-0.045115	-0.054337	-0.024360	-0.035905	

	deep_minutes	light_minutes	rem_minutes	\
dateOfSleep	-0.116257	0.128219	-0.052281	
bpm_max	0.040020	-0.014968	0.066116	
bpm_min	-0.019692	-0.317837	0.059835	
bpm_avg	0.029081	-0.262923	-0.002994	
bpm_std	0.042664	0.106530	0.034241	
bpm_median	-0.025979	-0.369781	-0.082595	
bpm_10th	0.065534	-0.370895	0.068026	
bpm_90th	0.041224	0.007787	0.015588	
bpm_count	0.237387	0.200950	0.180984	
minutesAsleep	0.496881	0.682627	0.509290	
minutesAwake	0.025971	0.315229	-0.299799	
timeInBed	0.477823	0.715528	0.419898	
deep_count	0.359378	0.125919	0.169642	

light_count	0.250720	0.579264	0.077876
rem_count	0.292509	-0.089013	0.555045
wake_count	0.294964	0.486325	0.147443
deep_minutes	1.000000	-0.109591	0.329167
light_minutes	-0.109591	1.000000	-0.139736
rem_minutes	0.329167	-0.139736	1.000000
wake_minutes	0.025971	0.315229	-0.299799
deep_thirtyDayAvgMinutes	0.013402	-0.019801	0.006406
light_thirtyDayAvgMinutes	-0.113682	0.063728	-0.058717
rem_thirtyDayAvgMinutes	0.031306	-0.022027	0.009317
wake_thirtyDayAvgMinutes	0.005048	-0.026890	-0.037836
prev_hrv	-0.013861	0.212744	-0.065960
target_hrv	-0.078151	0.027930	-0.175154

	wake_minutes	deep_thirtyDayAvgMinutes	\
dateOfSleep	0.044509	-0.092159	
bpm_max	-0.101243	0.011244	
bpm_min	-0.149808	0.110658	
bpm_avg	-0.101802	0.155230	
bpm_std	0.008633	-0.009105	
bpm_median	-0.124105	0.186094	
bpm_10th	-0.137978	0.185833	
bpm_90th	-0.009024	0.036697	
bpm_count	0.135607	-0.002358	
minutesAsleep	0.118121	-0.006983	
minutesAwake	1.000000	0.028399	
timeInBed	0.328420	-0.002073	
deep_count	0.094187	0.077491	
light_count	0.260005	0.053393	
rem_count	-0.132725	0.057468	
wake_count	0.166249	0.092272	
deep_minutes	0.025971	0.013402	
light_minutes	0.315229	-0.019801	
rem_minutes	-0.299799	0.006406	
wake_minutes	1.000000	0.028399	
deep_thirtyDayAvgMinutes	0.028399	1.000000	
light_thirtyDayAvgMinutes	0.012103	0.001759	
rem_thirtyDayAvgMinutes	-0.024911	0.445296	
wake_thirtyDayAvgMinutes	0.054628	0.375852	
prev_hrv	0.099092	-0.131735	
target_hrv	0.060423	-0.104745	

	light_thirtyDayAvgMinutes	rem_thirtyDayAvgMinutes	\
dateOfSleep	0.222209	-0.046121	
bpm_max	-0.048390	-0.021180	
bpm_min	-0.204543	0.031532	
bpm_avg	-0.219024	0.031240	

bpm_std	0.015408	-0.009556
bpm_median	-0.243644	0.029352
bpm_10th	-0.296932	0.056902
bpm_90th	-0.057919	-0.002272
bpm_count	-0.020538	-0.073252
minutesAsleep	-0.025883	0.000271
minutesAwake	0.012103	-0.024911
timeInBed	-0.017975	0.002071
deep_count	-0.052548	0.000782
light_count	-0.003848	-0.024750
rem_count	-0.099303	0.035218
wake_count	-0.025634	0.023740
deep_minutes	-0.113682	0.031306
light_minutes	0.063728	-0.022027
rem_minutes	-0.058717	0.009317
wake_minutes	0.012103	-0.024911
deep_thirtyDayAvgMinutes	0.001759	0.445296
light_thirtyDayAvgMinutes	1.000000	0.001085
rem_thirtyDayAvgMinutes	0.001085	1.000000
wake_thirtyDayAvgMinutes	0.409954	-0.085060
prev_hrv	0.189455	-0.076416
target_hrv	0.166764	0.011357

	wake_thirtyDayAvgMinutes	prev_hrv	target_hrv
dateOfSleep	0.186204	0.025463	0.025688
bpm_max	-0.026494	0.027929	-0.294057
bpm_min	-0.006736	-0.729785	-0.225867
bpm_avg	-0.053862	-0.402024	-0.318599
bpm_std	-0.036509	0.181155	-0.208096
bpm_median	-0.037234	-0.393489	-0.217225
bpm_10th	-0.028147	-0.810916	-0.237110
bpm_90th	-0.051236	-0.073119	-0.235557
bpm_count	0.011239	0.005807	-0.264832
minutesAsleep	-0.037440	0.116860	-0.101370
minutesAwake	0.054628	0.089138	0.052049
timeInBed	-0.028304	0.128224	-0.083061
deep_count	0.040581	0.128256	-0.045115
light_count	-0.056804	0.139161	-0.054337
rem_count	-0.101375	-0.128832	-0.024360
wake_count	-0.070096	0.083330	-0.035905
deep_minutes	0.005048	-0.013861	-0.078151
light_minutes	-0.026890	0.212744	0.027930
rem_minutes	-0.037836	-0.065960	-0.175154
wake_minutes	0.054628	0.099092	0.060423
deep_thirtyDayAvgMinutes	0.375852	-0.131735	-0.104745
light_thirtyDayAvgMinutes	0.409954	0.189455	0.166764
rem_thirtyDayAvgMinutes	-0.085060	-0.076416	0.011357

wake_thirtyDayAvgMinutes	1.000000	-0.027819	-0.035322
prev_hrv	-0.027819	1.000000	0.232980
target_hrv	-0.035322	0.232980	1.000000


```
[24]: import numpy as np
corr = final_df.corr()

# Generate a mask for the upper triangle
mask = np.triu(np.ones_like(corr, dtype=bool))

# Set up the matplotlib figure
f, ax = plt.subplots(figsize=(11, 9))

# Generate a custom diverging colormap
cmap = sns.diverging_palette(230, 20, as_cmap=True)

# Draw the heatmap with the mask and correct aspect ratio
sns.heatmap(corr, mask=mask, cmap=cmap, vmax=0.9, vmin=-0.9, center=0,
            square=True, linewidths=.5, cbar_kws={"shrink": .5})
plt.title('Correlation of Variables')
plt.show()
```

Based on what I can see we have high (≥ 0.90) correlation between a few different variables **wake_minutes** & **minutesAwake**, **timeInBed** & **minutesAsleep**, and **bpm_std** & **bpm_90th**. So I've decided to drop the **bpm_std**, **minutesAwake**, & **timeInBed** so the final filtered dataframe is below.

```
[25]: final_filt_df = final_df[['dateOfSleep', 'bpm_max', 'bpm_min', 'bpm_avg',
    'bpm_median', 'bpm_10th', 'bpm_90th', 'bpm_count',
    'minutesAsleep', 'deep_count',
    'light_count', 'rem_count', 'wake_count', 'deep_minutes',
    'light_minutes', 'rem_minutes', 'wake_minutes',
    'deep_thirtyDayAvgMinutes', 'light_thirtyDayAvgMinutes',
    'rem_thirtyDayAvgMinutes', 'wake_thirtyDayAvgMinutes', 'prev_hrv',
    'target_hrv']]
```

1.8 Conclusion

So we were able to use DuckDB to read in and manipulate json data, csv data, and dataframes with ease. We took the data we had with a goal in mind of being able to predict heart rate variability so we setup the data to do that and even filtered out highly correlated variables. Our data is now prepared for a machine learning model. Below I have put only the necessary code together to import and organize the data. I could probably make it more efficient but it runs in ~7 seconds to read in the sleep data, heart rate data, and heart rate variability data and manipulate it all to create our final dataset.

```
[26]: # organize it all
sleep_data = duckdb.read_json("./data/unzipped/Takeout/Fitbit/Global Export_
↳Data/sleep-*.json", timestamp_format="%Y-%m-%dT%H:%M:%S.%g")
sleep_meta_df = duckdb.sql("""Select distinct logId, dateOfSleep, startTime,
↳endTime, duration, minutesToFallAsleep, minutesAsleep,
minutesAwake, minutesAfterWakeup, timeInBed,
↳efficiency, type, infoCode, logType, mainSleep
FROM sleep_data
WHERE mainSleep == true""").df()
sleep_summary_df = duckdb.sql("""with summary_data as (select logId,
↳unnest(levels.summary, max_depth := 1) from sleep_data),
unioned_data as (
select logId, 'deep' as stage, unnest(deep) from summary_data
union all
select logId, 'wake' as stage, unnest(wake) from summary_data
union all
select logId, 'light' as stage, unnest(light) from summary_data
union all
select logId, 'rem' as stage, unnest(rem) from summary_data)
select logId, stage, count, minutes, max(thirtyDayAvgMinutes) as
↳thirtyDayAvgMinutes
from unioned_data
group by logId, stage, count, minutes
order by logId, stage""").df()
sleep_data_detail = duckdb.sql("""select logId, unnest(levels.data, recursive :
↳= True) from sleep_data order by logId""")
sleep_short_data = duckdb.sql("""select logId, unnest(levels.shortData,
↳recursive := True) from sleep_data order by logId""")
heart_rate_data = duckdb.read_json('./data/unzipped/Takeout/Fitbit/Global_
↳Export Data/heart_rate-*.json', timestamp_format="%m/%d/%y %H:%M:%S")
heart_rate_df = duckdb.sql('select dateTime, unnest(value) from
↳heart_rate_data').df()
hrv_df = duckdb.read_csv('./data/unzipped/Takeout/Fitbit/Heart Rate Variability/
↳Daily Heart Rate Variability Summary*.csv', timestamp_format="%Y-%m-%dT%H:%M:
↳%S").df()
adj_sleep_summary_df = duckdb.sql("""pivot sleep_summary_df on stage
using
```

```

        first(count) as count,
        first(minutes) as minutes,
        first(thirtyDayAvgMinutes) as thirtyDayAvgMinutes
    order by logId""").df()
complete_sleep_data = sleep_meta_df.merge(adj_sleep_summary_df, on='logId').
    ↪sort_values('dateOfSleep')
final_df = duckdb.sql("""WITH Complete_Sleep AS (
        select *,
            (dateOfSleep::Date + 1) as hrv_prediction_date,
            (startTime - Interval '2 Hours') as_
    ↪startTime_sleep_less_2_hrs,
            ((lead(startTime) over (order by logId)) - Interval '2_
    ↪Hours') as tomorrow_startTime_sleep_less_2_hrs
        from complete_sleep_data),
    HR_with_sleep AS (
        Select hr.*, cs.* from Complete_Sleep cs, heart_rate_df hr
        where hr.dateTime >= cs.startTime_sleep_less_2_hrs
        and hr.dateTime < tomorrow_startTime_sleep_less_2_hrs),
    hr_sleep_agg AS (
        Select dateOfSleep::date as dateOfSleep, hrv_prediction_date,
            max(bpm) as bpm_max, min(bpm) as bpm_min, round(avg(bpm), 2) as_
    ↪bpm_avg, round(stddev_samp(bpm), 2) as bpm_std,
            round(median(bpm),2) as bpm_median, round(quantile_cont(bpm, 0.
    ↪10),2) as bpm_10th, round(quantile_cont(bpm, 0.90),2) as bpm_90th,
            count(bpm) as bpm_count, minutesAsleep,
            minutesAwake,
            timeInBed,
            deep_count,
            light_count,
            rem_count,
            wake_count,
            deep_minutes,
            light_minutes,
            rem_minutes,
            wake_minutes,
            deep_thirtyDayAvgMinutes,
            light_thirtyDayAvgMinutes,
            rem_thirtyDayAvgMinutes,
            wake_thirtyDayAvgMinutes
        from HR_with_sleep
        Group by dateOfSleep::date, hrv_prediction_date, minutesAsleep,
    ↪minutesAwake,
            timeInBed,
            deep_count,
            light_count,
            rem_count,

```

```

        wake_count,
        deep_minutes,
        light_minutes,
        rem_minutes,
        wake_minutes,
        deep_thirtyDayAvgMinutes,
        light_thirtyDayAvgMinutes,
        rem_thirtyDayAvgMinutes,
        wake_thirtyDayAvgMinutes
    )
    select agg.*, lag(rmssd) over (order by dateOfSleep) as prev_hrv, hrv.
    ↳rmssd as target_hrv from hr_sleep_agg agg inner join hrv_df hrv on agg.
    ↳hrv_prediction_date = hrv.timestamp Order by dateOfSleep
    """).df()

final_filt_df = final_df[['dateOfSleep', 'bpm_max', 'bpm_min', 'bpm_avg',
    'bpm_median', 'bpm_10th', 'bpm_90th', 'bpm_count',
    'minutesAsleep', 'deep_count',
    'light_count', 'rem_count', 'wake_count', 'deep_minutes',
    'light_minutes', 'rem_minutes', 'wake_minutes',
    'deep_thirtyDayAvgMinutes', 'light_thirtyDayAvgMinutes',
    'rem_thirtyDayAvgMinutes', 'wake_thirtyDayAvgMinutes', 'prev_hrv',
    ↳'target_hrv']]

```